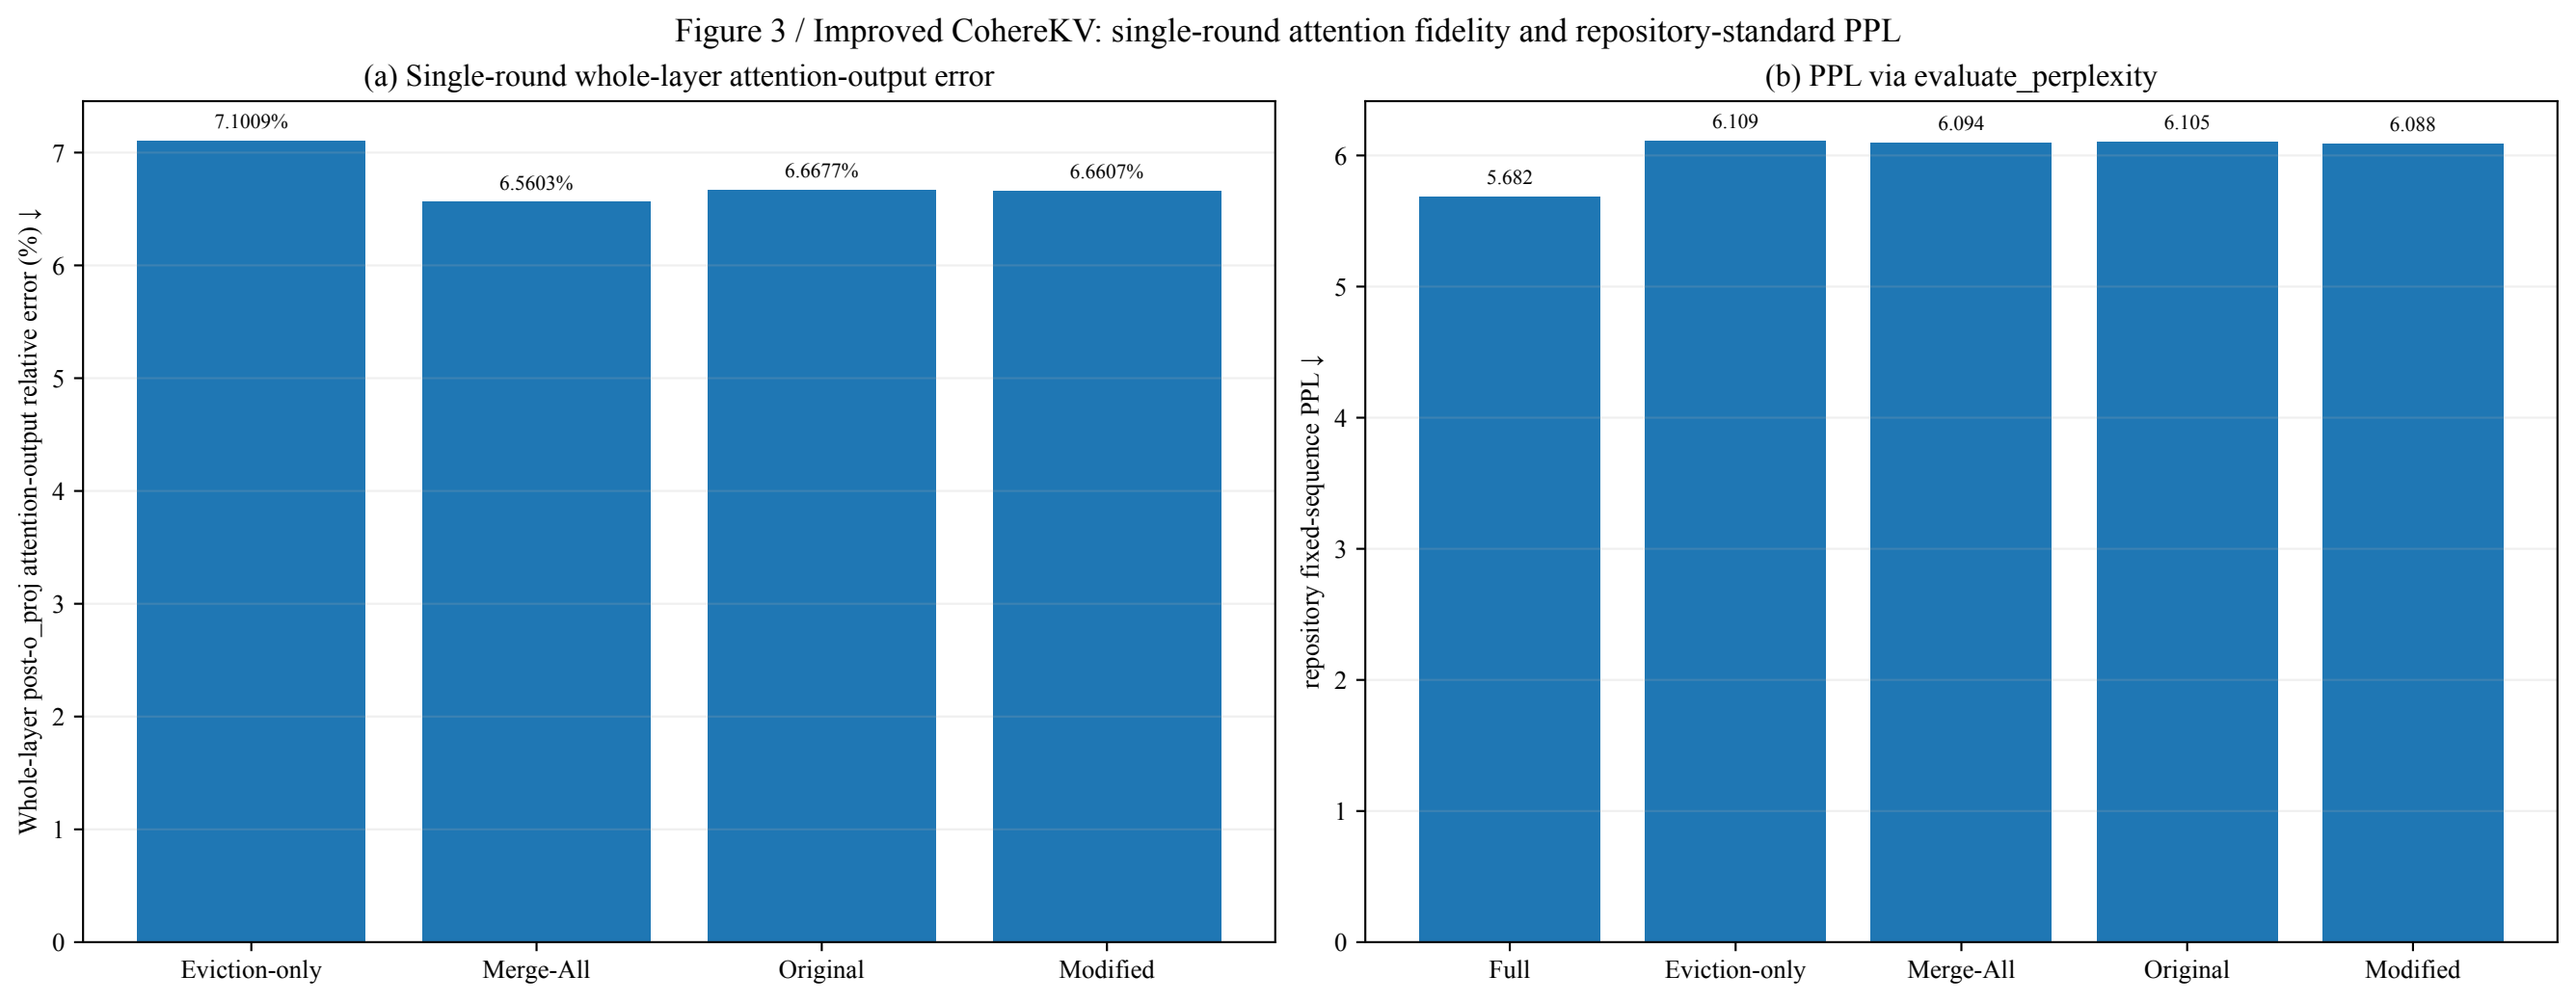
Figure 3 / Improved CohereKV: single-round attention fidelity and repository-standard PPL
(a) Single-round whole-layer attention-output error
Whole-layer post-o_proj attention-output relative error (%) ↓
7.1009%
6.5603%
6.6677%
6.6607%
Eviction-only
Merge-All
Original
Modified
(b) PPL via evaluate_perplexity
repository fixed-sequence PPL ↓
5.682
6.109
6.094
6.105
6.088
Full
Eviction-only
Merge-All
Original
Modified