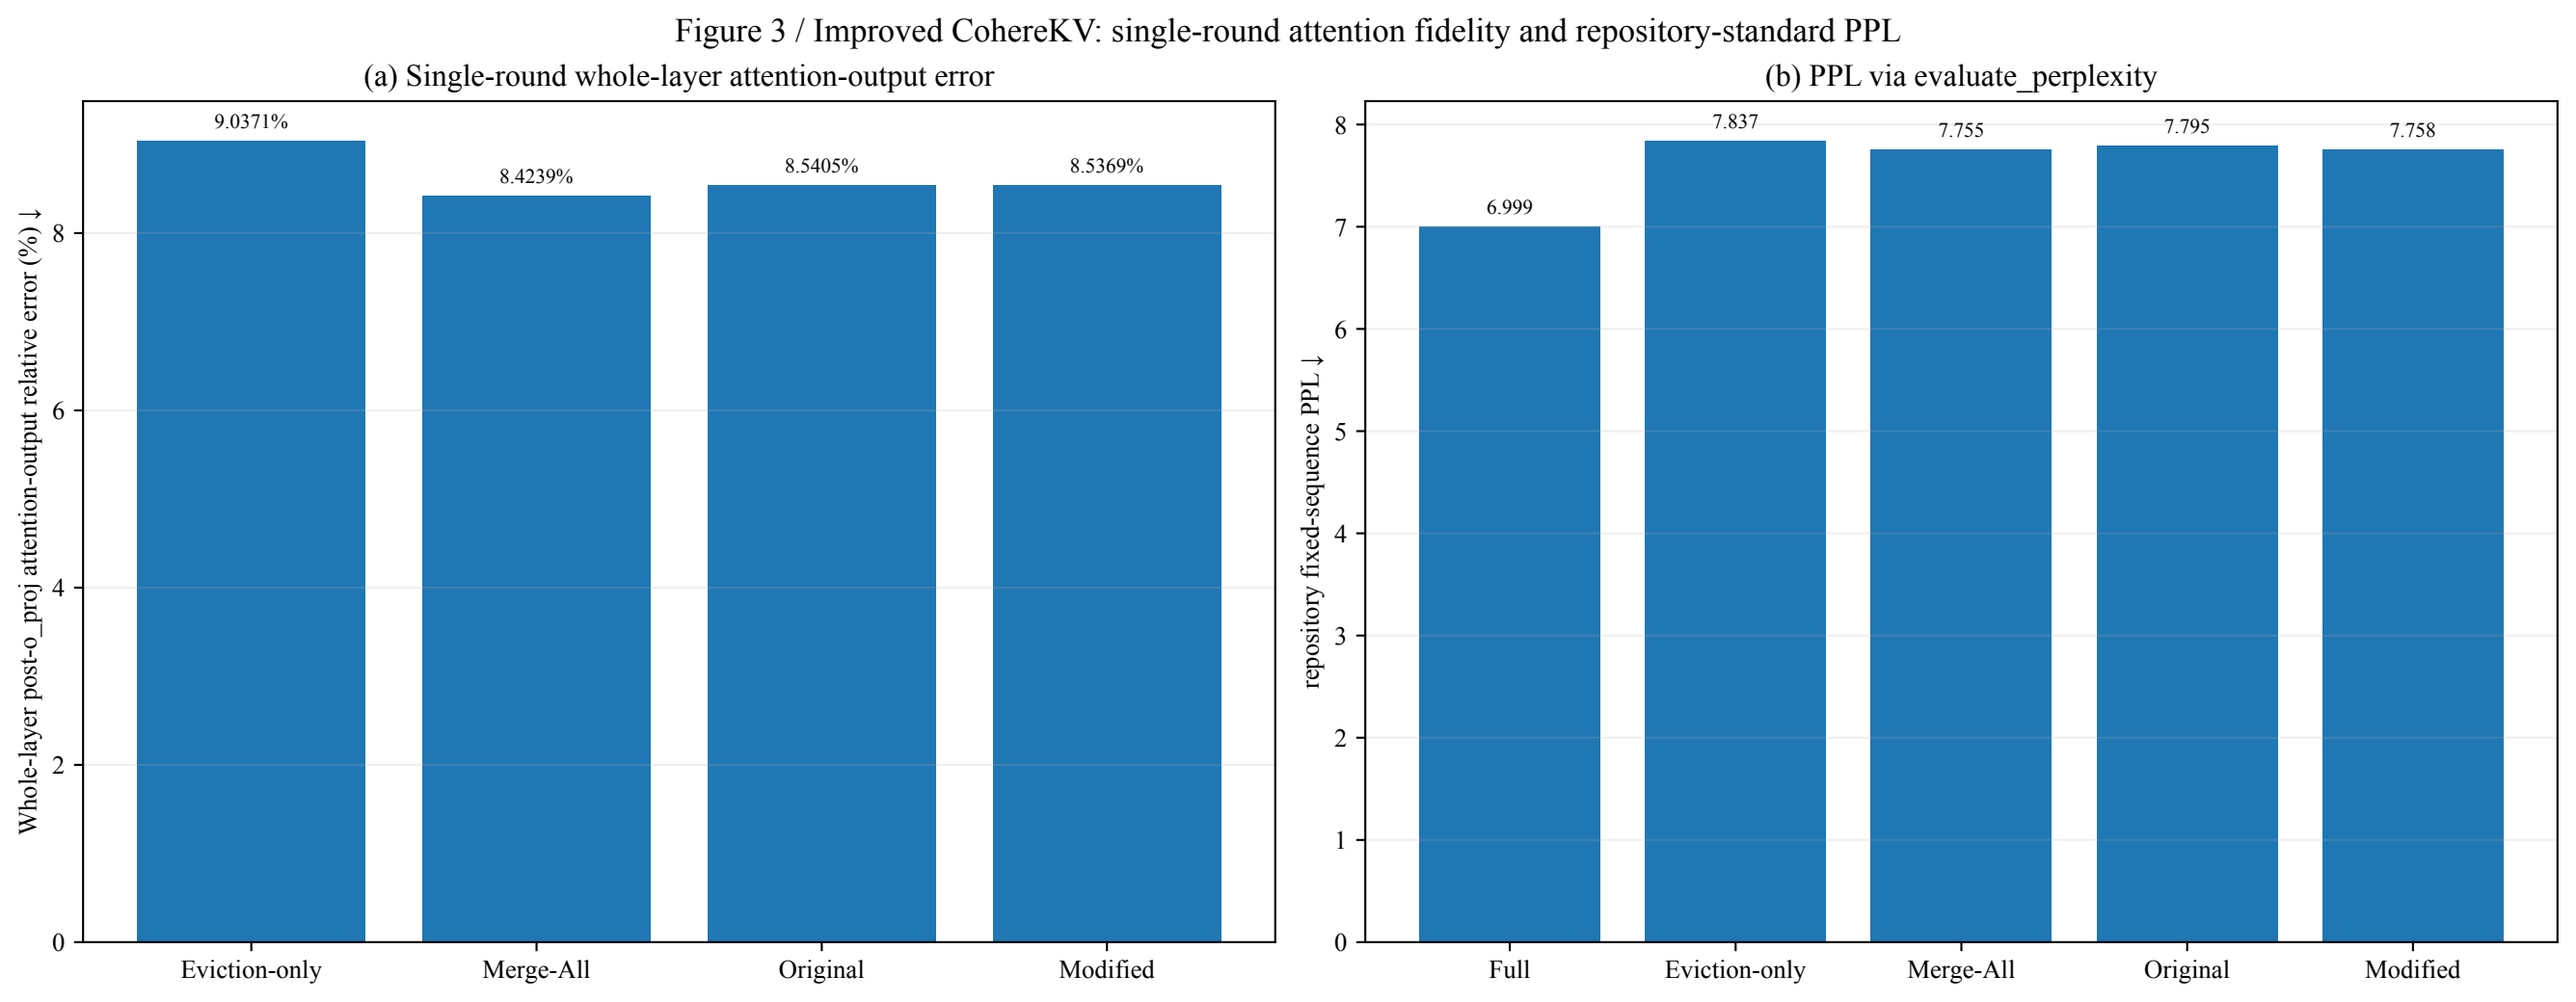
Figure 3 / Improved CohereKV: single-round attention fidelity and repository-standard PPL
(a) Single-round whole-layer attention-output error
Whole-layer post-o_proj attention-output relative error (%) ↓
9.0371%
8.4239%
8.5405%
8.5369%
Eviction-only
Merge-All
Original
Modified
(b) PPL via evaluate_perplexity
repository fixed-sequence PPL ↓
6.999
7.837
7.755
7.795
7.758
Full
Eviction-only
Merge-All
Original
Modified